

## ▼ California Housing Price Prediction :

### Problem Statement :

The US Census Bureau has published California Census Data which has 10 types of metrics such as the population, median income, median housing price, and so on for each block group in California. The dataset also serves as an input for project scoping and tries to specify the functional and nonfunctional requirements for it.

### Objective :

The project aims at building a model of housing prices to predict median house values in California using the provided dataset. This model should learn from the data and be able to predict the median housing price in any district, given all the other metrics.

Districts or block groups are the smallest geographical units for which the US Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people). There are 20,640 districts in the project dataset.

## ▼ Domain: Finance and Housing

### Analysis Tasks to be performed:

1. Build a model of housing prices to predict median house values in California using the provided dataset.
2. Train the model to learn from the data to predict the median housing price in any district, given all the other metrics.
3. Predict housing prices based on median\_income and plot the regression chart for it.

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
from pandas.plotting import scatter_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

%matplotlib inline
```

## ▼ 1. Load the data :

Read the "housing.xlsx" file from the folder into the program. Print first few rows of this data.  
Extract input (X) and output (Y) data from the dataset.

```
# Load the data using read_excel method in pandas
# housing_dt = pd.pandas.read_excel(r'/Users/arvindatmuri/PythonProjects/California
housing_dt = pd.pandas.read_excel(r'/content/california_housing_dataset.xlsx')
housing_dt.head(10)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populatio
0	-122.23	37.88	41	880		129.0
1	-122.22	37.86	21	7099		1106.0
2	-122.24	37.85	52	1467		190.0
3	-122.25	37.85	52	1274		235.0
4	-122.25	37.85	52	1627		280.0
5	-122.25	37.85	52	919		213.0
6	-122.25	37.84	52	2535		489.0
7	-122.25	37.84	52	3104		687.0
8	-122.26	37.84	42	2555		665.0
9	-122.25	37.84	52	3549		707.0

```
# Dataset Description using shape method()
print("Rows:", housing_dt.shape[0])
print("Columns:", housing_dt.shape[1])
```

```
Rows: 20640
Columns: 10
```

## ▼ Dataset Description :

Dataset Size : 20640 rows x 10 columns

```
# Count and Column Data Type
housing_dt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  int64
3   total_rooms           20640 non-null  int64
```

```

4   total_bedrooms      20433 non-null    float64
5   population          20640 non-null    int64
6   households          20640 non-null    int64
7   median_income       20640 non-null    float64
8   ocean_proximity     20640 non-null    object
9   median_house_value  20640 non-null    int64
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB

```

## ▼ Field Description

longitude (signed numeric - float) : Longitude value for the block in California, USA

latitude (numeric - float) : Latitude value for the block in California, USA

housing\_median\_age (numeric - int) : Median age of the house in the block

total\_rooms (numeric - int) : Count of the total number of rooms (excluding bedrooms) in all houses in the block

total\_bedrooms (numeric - float) : Count of the total number of bedrooms in all houses in the block

population (numeric - int) : Count of the total number of population in the block

households (numeric - int) : Count of the total number of households in the block

median\_income (numeric - float) : Median of the total household income of all the houses in the block

ocean\_proximity (numeric - categorical) : Type of the landscape of the block [ Unique Values : 'NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND' ]

median\_house\_value (numeric - int) : Median of the household prices of all the houses in the block

```

# Calculate all the Measures of Central Tendency
housing_dt.describe()

```

## ▼ 2. Handle missing values :

Fill the missing values with the mean of the respective column.

```
housing_dt.isnull().sum()
```

```
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms    207
population         0
households         0
median_income      0
ocean_proximity    0
median_house_value 0
dtype: int64
```

```
# Calculate all the Measures of Central Tendency(Mean, Median and Mode) for Total B
mean_total_bedrooms = housing_dt['total_bedrooms'].mean()
median_total_bedrooms = housing_dt['total_bedrooms'].median()
mode_total_bedrooms = housing_dt['total_bedrooms'].mode()
```

```
print("Mean:", mean_total_bedrooms)
print("Median: ", median_total_bedrooms)
# print("Mode:", mode_total_bedrooms)
print("Null Values: ", housing_dt['total_bedrooms'].isnull().sum())
```

```
Mean: 537.8705525375618
Median: 435.0
Null Values: 207
```

## ▼ Let's stick to Mean in this case, to replace the NA/Null values as per the task

```
# Filling Mean Values with Mean calculated above
housing_dt['total_bedrooms'].fillna(value = mean_total_bedrooms, inplace=True)
```

```
housing_dt.isnull().sum()
```

```
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms     0
population         0
households         0
median_income      0
ocean_proximity    0
median_house_value 0
dtype: int64
```

Our Second Task is also completed Here, We have filled all the Null values with its Mean

### ▼ 3. Encode categorical data :

Convert categorical column in the dataset to numerical data.

Looking at the data, all the columns are numerical except to ocean\_proximity Column. So lets convert the Categorical Data Column into Numerical Data.

```
label_encoder = LabelEncoder()
housing_dt['ocean_proximity'] = label_encoder.fit_transform(housing_dt['ocean_proxi
housing_dt['ocean_proximity'].sample(5)
```

12290	1
17276	0
10513	0
10940	0
5915	0

Name: ocean\_proximity, dtype: int64

### ▼ 4. Standardize data :

Standardize training and test datasets.

```
names=housing_dt.columns
st_sc = StandardScaler()

scaled_housing_dt = st_sc.fit_transform(housing_dt)
scaled_housing_dt = pd.DataFrame(scaled_housing_dt, columns=names)
scaled_housing_dt.sample(5)
```


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	pc
<b>13738</b>	1.177801	-0.740606	-1.957806	0.122956	0.536647	
<b>16406</b>	-0.803748	1.104049	0.187562	-0.085151	0.050144	
<b>16724</b>	-0.544200	-0.075781	-0.845393	-0.074150	-0.210906	
<b>3985</b>	0.464044	-0.675060	0.346478	0.174754	-0.104113	
<b>15548</b>	1.237697	-1.176020	-1.401600	-0.948294	-0.839800	

### ▼ 5. Visualize the Data to Check the linearity among columns :

Lets put the preprocessed data into a graph to identify the linearity relationship between any two columns. To Visualize the Data Lets use the scaled data for better results.


```
sns.pairplot(scaled_housing_dt, markers=["o"])
```

<seaborn.axisgrid.PairGrid at 0x7fdfe3073b50>



```
# Heatmap using seaborn
housing_corr_matrix = scaled_housing_dt.corr()
fig, axe = plt.subplots(figsize=(10,8))
cmap = sns.diverging_palette(250,10, center = "light", as_cmap=True)
sns.heatmap(housing_corr_matrix,square =True, cmap=cmap, annot=True )
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fdfe0ec0790>



Shades and Density of Red Color shows the stronger relation between columns.

lon lat hmae tot\_rooms tot\_bedrooms pop households med\_income ocean\_prox med\_house\_value

From the Above Diagram/Plots, These are be our observations:

1. Linearity can be observed among very less columns, Not all can be considered for our analysis.
2. Latitude vs Longitude - Analysing a relationship, and hence Doesn't make much sense since data provided for these two columns are just Co-ordinates
3. Total\_bedrooms vs Total Rooms - Analysing a relationship between these columns also doesn't make much sense
4. Households vs Total Bedrooms, Total rooms, Population - Also has linearity between them, but this may not be useful for our problem statement.
5. Last but not least, there is a linearity observed with the Median income and Median House Value - which has a strong relationship with income and house price and should be considered for our problem statement. Also as per the heatmap, Only column very much relatable to the output column is median\_income)

According to Dataset and Problem Statement, we have to find out the Median House Value so this column is considered as Dependent Column/Target Column and also can be classified as Output column. Rest of the columns can be classified as Independent Data.

```
scaled_housing_dt.plot.scatter(x = 'median_income', y = 'median_house_value')
```



\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be `<matplotlib.axes._subplots.AxesSubplot at 0x7fdfe12d2a50>`



## ▼ 6. Split the dataset :

Split the data into 80% training dataset and 20% test dataset. Using `train_test_split` to split the dataset into train and test methods.

```
X = scaled_housing_dt.drop('median_house_value',axis=1)
Y = scaled_housing_dt["median_house_value"]
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2,random_stat
x_train.shape, y_train.shape, x_test.shape, y_test.shape

((16512, 9), (16512,), (4128, 9), (4128,))
```

## ▼ 7. Perform Linear Regression :

[ HINT: Import `mean_squared_error` from `sklearn.metrics` ]

a. Perform Linear Regression on training data.

```
lr = LinearRegression()
lr.fit(x_train,y_train)

LinearRegression()
```

b. Predict output for test dataset using the fitted model.

```
import math
y_pred = lr.predict(x_test)
```

c. Print root mean squared error (RMSE) from Linear Regression.

```
print("Model Score", lr.score(x_train, y_train))
print("RMSE:", math.sqrt(mean_squared_error(y_test, y_pred)))
print("R2 Score:", r2_score(y_test, y_pred))
```

```
Model Score 0.635296662157909
RMSE: 0.5998770880694679
R2 Score: 0.6366841281048751
```

```
lm = smf.ols(formula='median_house_value ~ longitude+latitude+housing_median_age+total_rooms+total_bedrooms+population+households+median_income+ocean_proximity')
lm.summary()
```

↗

OLS Regression Results						
<b>Dep. Variable:</b>	median_house_value	<b>R-squared:</b>	0.636			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.635			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3999.			
<b>Date:</b>	Tue, 07 Dec 2021	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	13:00:16	<b>Log-Likelihood:</b>	-18868.			
<b>No. Observations:</b>	20640	<b>AIC:</b>	3.776e+04			
<b>Df Residuals:</b>	20630	<b>BIC:</b>	3.783e+04			
<b>Df Model:</b>	9					
<b>Covariance Type:</b> nonrobust						
	coef	std err	t	P> t	[0.025	0.975]
Intercept	-3.469e-17	0.004	-8.26e-15	1.000	-0.008	0.008
longitude	-0.7393	0.013	-57.263	0.000	-0.765	-0.714
latitude	-0.7858	0.013	-61.664	0.000	-0.811	-0.761
housing_median_age	0.1248	0.005	26.447	0.000	0.116	0.134
total_rooms	-0.1265	0.015	-8.609	0.000	-0.155	-0.098
total_bedrooms	0.2995	0.022	13.630	0.000	0.256	0.343
population	-0.3907	0.011	-36.927	0.000	-0.411	-0.370
households	0.2589	0.022	11.515	0.000	0.215	0.303
median_income	0.6549	0.005	119.287	0.000	0.644	0.666
ocean_proximity	0.0009	0.005	0.190	0.850	-0.008	0.010
<b>Omnibus:</b>	5037.491	<b>Durbin-Watson:</b>	0.965			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	18953.000			
<b>Skew:</b>	1.184	<b>Prob(JB):</b>	0.00			
<b>Kurtosis:</b>	7.054	<b>Cond. No.</b>	14.2			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## ▼ 7. Bonus exercise: Perform Linear Regression with one independent variable :

Extract just the median\_income column from the independent variables (from X\_train and X\_test).

```
x_train_in = x_train[["median_income"]]
x_test_in = x_test[["median_income"]]
x_train_in.shape, x_test_in.shape
```

```
((16512, 1), (4128, 1))
```

Perform Linear Regression to predict housing values based on median\_income.

```
lr1 = LinearRegression()
lr1.fit(x_train_in, y_train)
```

```
LinearRegression()
```

Predict output for test dataset using the fitted model.

```
y_pred_in = lr1.predict(x_test_in)
```

Print root mean squared error (RMSE) from Linear Regression.

```
print("Model Score", lr1.score(x_train_in, y_train))
print("RMSE:", math.sqrt(mean_squared_error(y_test, y_pred_in)))
print("R2 Score:", r2_score(y_test, y_pred_in))
```

```
Model Score 0.4723426637102913
RMSE: 0.7191095397723252
R2 Score: 0.4779045336215766
```

```
lm = smf.ols(formula='median_house_value ~ median_income', data=scaled_housing_dt).
lm.summary()
```

```

                OLS Regression Results
    Dep. Variable:  median_house_value    R-squared:  0.473
      Model:        OLS                   Adj. R-squared: 0.473
    Method:        Least Squares         F-statistic:  1.856e+04
      Date:         Tue, 07 Dec 2021      Prob (F-statistic): 0.00
      Time:         13:00:28              Log-Likelihood: -22668.
No. Observations: 20640                  AIC:         4.534e+04
Df Residuals:     20638                  BIC:         4.536e+04
Df Model:         1
Covariance Type:  nonrobust

               coef    std err          t      P>|t| [0.025 0.975]
Intercept    1.735e-16  0.005   3.43e-14  1.000  -0.010  0.010
median_income 0.6881    0.005  136.223  0.000   0.678  0.698
Omnibus:      4245.795    Durbin-Watson:  0.655
Prob(Omnibus): 0.000    Jarque-Bera (JB): 9273.446
Skew:         1.191     Prob(JB):      0.00
Kurtosis:     5.260     Cond. No.    1.00


```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot the fitted model for training data as well as for test data to check if the fitted model satisfies the test data.

```
scaled_housing_dt.plot(kind='scatter', x='median_income', y='median_house_value')  
plt.plot(x_test_in, y_pred_in, c="red", linewidth=1)  
  
*c* argument looks like a single numeric RGB or RGBA sequence, which should be  
[<matplotlib.lines.Line2D at 0x7fdfell1ebc10>]
```



✓ 0s completed at 02:09

● ✕